

Big Data Assignment

ISTD, SUTD

March 29, 2024

Group Assignment

This is a group assignment. You are supposed to work as a team based on your project group. Each group is only required to submit one copy of your solution through eDimension.

Deadline

14 April 2024 23:58.

Synopsis

In this assignment, you are supposed to develop a big data application to process user generated data capturing customer opinions towards brands and businesses.

Your tasks include loading the raw data into the Hadoop distributed file system (HDFS), performing data transformation and cleaning using Spark framework. Finally, you should be able to conduct some descriptive analytics on the cleaned data.

Common Requirements

For all the questions,

1. You may use RDD API and/or Dataframe API. You are not allowed to use Spark SQL API. e.g. `spark.sql("SELECT")` will not be rewarded with any mark.
2. The answer should not be dependent on the results from other questions, i.e. for each answer script, it should start from reading the given raw data file. However, you may develop your own library to be shared between answers for different questions.
3. There should be only *one* call to `read()` and *one* call to `write()` in each answer script.

4. Place the input data in the designated HDFS path and output data should be written in the designated HDFS path.

Submission

You are supposed to submit the following

1. q1.py
2. q2.py
3. a2.py
4. q4.py
5. q5.py
6. hw2.sh

Each py file should contain the solution to the correspondent question, i.e. q1.py for question 1, q2.py for question 2 so on and so forth. The script hw2.sh sets up the needed folders and data in HDFS and submits the python scripts to the spark cluster.

Templates of these six files are given to you. Besides editing the python template files, please read through the hw2.sh file and update the following sections

```
# change the following according to your student numbers  
echo "1001234,1003456"
```

```
# change the following according to your environment  
hdfs_namenode="localhost"
```

You may assume that during the grading process, the data folder ./data, the .py files and hw2.sh file are placed under the same linux folder.

Part 1

Data

In part 1 we are looking at the restaurant review data extracted from

<https://www.kaggle.com/damienbeneschi/krakow-ta-restaurants-data-raw>

You **don't** need to and are **not recommended** to download the data from kaggle. Please make use of the data provided to you along with this assignment.

Copy the CSV file TA_restaurants_curated_cleaned.csv into HDFS path /assignment2/part1/input/. Load it into a Dataframe or RDD, we may observe that the data set has the following schema

```
hdfs_nn = "localhost"  
df = spark.read.option("header",True)\
```

```
.csv("hdfs://s:9000/assignment2/part1/input/" % (hdfs_nn))
df.printSchema()
```

```
root
 |-- _c0: string (nullable = true)
 |-- Name: string (nullable = true)
 |-- City: string (nullable = true)
 |-- Cuisine Style: string (nullable = true)
 |-- Ranking: string (nullable = true)
 |-- Rating: string (nullable = true)
 |-- Price Range: string (nullable = true)
 |-- Number of Reviews: string (nullable = true)
 |-- Reviews: string (nullable = true)
 |-- URL_TA: string (nullable = true)
 |-- ID_TA: string (nullable = true)
```

and a preview of the data records looks like the following:

_c0	Name	City	Cuisine Style	Ranking	Rating	Price Range	Number of Reviews	Reviews	URL_TA	ID_TA
0	Martine of Martin...	Amsterdam	['French', 'Dutch...	1.0	5.0	\$ - \$\$\$	136.0	['Just like home...	/Restaurant_Revie...	d11752080
1	De Silveren Spiegel	Amsterdam	['Dutch', 'Europe...	2.0	4.5	\$\$\$\$	812.0	['Great food and...	/Restaurant_Revie...	d693419
2	La Rive	Amsterdam	['Mediterranean',...	3.0	4.5	\$\$\$\$	567.0	['Satisfaction',...	/Restaurant_Revie...	d696959
3	Vinkeles	Amsterdam	['French', 'Europ...	4.0	5.0	\$\$\$\$	564.0	['True five star...	/Restaurant_Revie...	d1239229
4	Librije's Zusje A...	Amsterdam	['Dutch', 'Europe...	5.0	4.5	\$\$\$\$	316.0	['Best meal....	/Restaurant_Revie...	d6864170
5	Ciel Bleu Restaurant	Amsterdam	['Contemporary', ...	6.0	4.5	\$\$\$\$	745.0	['A treat!', 'Wo...	/Restaurant_Revie...	d696902
6	Zaza's	Amsterdam	['French', 'Inter...	7.0	4.5	\$ - \$\$\$	1455.0	['40th Birthday ...	/Restaurant_Revie...	d1014732
7	Blue Pepper Resta...	Amsterdam	['Asian', 'Indone...	8.0	4.5	\$\$\$\$	675.0	['Great Experien...	/Restaurant_Revie...	d697058
8	Teppanyaki Restau...	Amsterdam	['Japanese', 'Asi...	9.0	4.5	\$\$\$\$	923.0	['Great Food & S...	/Restaurant_Revie...	d697009
9	Rob Wigboldus Vis...	Amsterdam	['Dutch', 'Seafoo...	10.0	4.5	\$	450.0	['Excellent Herr...	/Restaurant_Revie...	d1956562
10	The Happy Bull	Amsterdam	['American', 'Bar...	11.0	4.5	\$ - \$\$\$	295.0	['Simply AMAZING...	/Restaurant_Revie...	d10275170
11	Gartine	Amsterdam	['French', 'Dutch...	12.0	4.5	\$ - \$\$\$	967.0	['A hidden gem',...	/Restaurant_Revie...	d1014753
12	Restaurant Adam	Amsterdam	['French', 'Europ...	13.0	4.5	\$\$\$\$	368.0	['Love it!', 'As...	/Restaurant_Revie...	d7695005
13	Biercafe Gollern	Amsterdam	['Bar', 'Pub']	14.0	4.5	\$ - \$\$\$	586.0	['Awesome little...	/Restaurant_Revie...	d3893242
14	Restaurant Duelder	Amsterdam	['French', 'Dutch...	15.0	4.5	\$\$\$\$	1246.0	['Best meal of o...	/Restaurant_Revie...	d1408533
15	Greenwoods Keizer...	Amsterdam	['Dutch', 'Cafe',...	16.0	4.5	\$ - \$\$\$	1391.0	['So. Much. Food...	/Restaurant_Revie...	d3200493
16	Omegg - City Ce...	Amsterdam	['Dutch', 'Europe...	17.0	4.5	\$	1633.0	['Brunch', 'Wor...	/Restaurant_Revie...	d8562698
17	Brasserie Ambassade	Amsterdam	['French', 'Bar',...	18.0	4.5	\$\$\$\$	958.0	['Wonderful Chri...	/Restaurant_Revie...	d8567150
18	Sherpa Restaurant	Amsterdam	['Indian', 'Tibet...	19.0	4.5	\$ - \$\$\$	426.0	['Very good tibe...	/Restaurant_Revie...	d602573
19	La Maschera Lillo...	Amsterdam	['Italian', 'Medi...	20.0	4.5	\$ - \$\$\$	421.0	['Fabulous Itali...	/Restaurant_Revie...	d10071792

only showing top 20 rows

Question 1 (1 mark)

Develop a Spark application that cleans up the CSV file by removing rows with no reviews or rating < 1.0. Write the output as **CSV** into HDFS path /assignment2/output/question1/.

Sample output

_c0	Name	City	Cuisine Style	Ranking	Rating	Price Range	Number of Reviews	Reviews	URL_TA	ID_TA
1700	Auberge de la Rei...	Paris	['French', 'Euro...	701.0	4.0	\$ - \$\$\$	489.0	['Cozy Restaur...	/Restaurant_Revie...	d695128
1701	Le Petit Vendome	Paris	['French', 'Euro...	702.0	4.0	\$ - \$\$\$	343.0	['Parisian way...	/Restaurant_Revie...	d1146488
1702	La Cave Lanrezac	Paris	['Wine Bar', 'Eu...	703.0	4.5	\$ - \$\$\$	178.0	['Dinner with ...	/Restaurant_Revie...	d812970
1704	Chez Fernand Chri...	Paris	['French', 'Euro...	705.0	4.0	\$ - \$\$\$	892.0	['Tourist Area...	/Restaurant_Revie...	d1580042

Question 2 (1 mark)

Develop a Spark application that finds the best and the worst restaurants for each city for each price range (in terms of rating). Write the output as **CSV**

into HDFS path `/assignment2/output/question2/`. For simplicity, you can ignore rows with `Price Range` field as `null`.

You may use RDD API and/or Dataframe API. You are not allowed to use Spark SQL API.

Sample output:

_c0	Name	City	Cuisine Style	Ranking	Rating	Price Range	Number of Reviews	Reviews	URL_TA	ID_TA
3198	Pietersma Snacks	Amsterdam	['Dutch', 'Europ...	3209.0	5.0	\$	null	[[], []]	/Restaurant_Review...	d10587448
2932	Grillroom Sabba	Amsterdam	['Middle Eastern']	2942.0	2.5	\$	12.0	[['This is a gr...	/Restaurant_Review...	d6464568
1503	1 Chefalyon	Lyon	['Pub', 'Gastrop...	1485.0	5.0	\$\$\$	null	[[], []]	/Restaurant_Review...	d12408653
2605	Papagayo	Lyon	['Diner']	2606.0	2.0	\$\$\$	33.0	[[], []]	/Restaurant_Review...	d1329792
2951	le bountje	Brussels	['Belgian', 'Eur...	2952.0	5.0	\$ - \$\$\$	null	[[], []]	/Restaurant_Review...	d1563747
3009	Belga & Co	Brussels	['European']	null	-1.0	\$ - \$\$\$	null	[[], []]	/Restaurant_Review...	d13531979
2462	Sushi Express	Stockholm	['Japanese', 'Su...	null	5.0	\$ - \$\$\$	2.0	[[], []]	/Restaurant_Review...	d13344590

Question 3 (1 mark)

Develop a Spark application that extracts the three cities with the highest and lowest average rating per restaurant. Combine them, sorted, such that the output looks like this:

For instance:

City	AverageRating	RatingGroup
Athens	4.241316931982634	Top
London	4.178003263308178	Top
Krakow	4.164012738853503	Top
Geneva	3.97270245677889	Bottom
Helsinki	3.9153318077803205	Bottom
Brussels	3.900580875781948	Bottom

Write the output as **CSV** files into HDFS path `/assignment2/output/question3/`.

Question 4 (1 mark)

Develop a Spark application that counts the number of restaurants by city and cuisine style.

The output should something like the following:

City	Cuisine	count
Amsterdam	Vietnamese	24
Bratislava	Hungarian	3
Brussels	International	74
London	Kosher	26
Lyon	Mediterranean	80

| Lyon | German | 2 |

Write the output as **CSV** files into HDFS path `/assignment2/output/question4/`.

Part 2

In this second part of the assignment, we investigate the movie credit data extracted from

https://www.kaggle.com/tmdb/tmdb-movie-metadata?select=tmdb_5000_credits.csv

You don't need to and are not recommended to download the data from kaggle. Please make use of the data provided to you along with this part.

For part 2, instead of CSV, we consider the input file in **Parquet** format. Parquet format is a compressed column based format which is optimized for parallel processing. For more details of parquet file format, refer to the following documentation:

<https://spark.apache.org/docs/latest/sql-data-sources-parquet.html>

Copy the `tmdb_5000_credits.parquet` file into HDFS path `/part2/input/`. Load it into a Dataframe or RDD, we may observe that the data set has the following schema

```
df = spark.read.option("header",True)\
    .parquet("hdfs://%s:9000/assignment2/part2/input/" % (hdfs_nn))
df.printSchema()

root
 |-- movie_id: long (nullable = true)
 |-- title: string (nullable = true)
 |-- cast: string (nullable = true)
 |-- crew: string (nullable = true)
```

If we take a look at the first rows of the data, we see the following:

```
+-----+-----+-----+-----+
|movie_id|title|cast|crew|
+-----+-----+-----+-----+
| 19995|Avatar|[{cast_id: 242,...}|{credit_id: "5...|
| 285|Pirates of the Ca...|[{cast_id: 4, "...}|{credit_id: "5...|
| 206647|Spectre|[{cast_id: 1, "...}|{credit_id: "5...|
| 49026|The Dark Knight R...|[{cast_id: 2, "...}|{credit_id: "5...|
| 49529|John Carter|[{cast_id: 5, "...}|{credit_id: "5...|
| 559|Spider-Man 3|[{cast_id: 30, ...}|{credit_id: "5...|
| 38757|Tangled|[{cast_id: 34, ...}|{credit_id: "5...|
| 99861|Avengers: Age of ...|[{cast_id: 76, ...}|{credit_id: "5...|
| 767|Harry Potter and ...|[{cast_id: 3, "...}|{credit_id: "5...|
```

Question 5 (2 marks)

Develop a Spark application that finds the pairs of actors/actresses that are co-cast for **at least 2** movies. The output should be in a (set of) **Parquet** files in the following schema:

```
movie_id, title, actor1, actor2
```

Note that the result should not contain any repetition, e.g.

```
49026, The Dark Knight Rises, Michael Caine, Christian Bale
```

is considered as a duplicate entry of

```
49026, The Dark Knight Rises, Christian Bale, Michael Caine
```

The output should be something like the following:

movie_id	title	actor1	actor2
69848	One Man's Hero	James Gammon	Tom Berenger
9942	Major League	James Gammon	Tom Berenger
285	Pirates of the Ca...	David Bailie	Ho-Kwan Tse
58	Pirates of the Ca...	David Bailie	Ho-Kwan Tse
921	Cinderella Man	Michael Stevens	Conrad Bergschneider
14577	Dirty Work	Michael Stevens	Conrad Bergschneider
16290	Jackass 3D	Dimitry Elyashkevich	Manny Puig
12094	Jackass Number Two	Dimitry Elyashkevich	Manny Puig
9012	Jackass: The Movie	Dimitry Elyashkevich	Manny Puig

Hint

You should be able to extract the needed info from the `movie_id`, `title` and `cast` columns.

Good luck!